


Embracing the

Curiosity Economy



How human judgment, experience,
and curiosity become your greatest
assets in the age of AI.

David Fearne

Written with AI, Born from Curiosity

There is something I should tell you before you read this book.

I am dyslexic.

That sentence probably seems relatively unremarkable. Millions of people are dyslexic. Many of them are extraordinary writers. Dyslexia does not prevent someone from writing a book, and I certainly don't want to suggest that it does.

But it matters to the story of *this* book.

In fact, this is probably the first time in my professional life that I have openly told people that I am dyslexic.

That was a conscious choice.

Throughout my career, I never wanted to be judged through the lens of dyslexia. I didn't want somebody looking at a piece of work differently because they knew. I didn't want an achievement qualified by it. I didn't want someone deciding what I might or might not be capable of before I had the opportunity to show them.

I wanted to be judged by my work.

Whether that was the right decision is something I still don't entirely know. But it was my decision.

And so I adapted.

I built a career in technology. I became an AI leader. I advised organisations and governments, built systems, led teams, stood on stages, developed ideas and spent an extraordinary amount of my life asking questions about technology and what it might mean for the people who eventually have to live with it.

I have never struggled for ideas.

Writing those ideas down has been different.

I can talk for hours about something that interests me. I can take an idea apart, challenge it, connect it to something seemingly unrelated, rebuild it and then disappear down another intellectual rabbit hole because of something I have just realised along the way.

Give me a question and I will explore it.

Give me a blank page and tell me to write 4,000 words about it, and that is an entirely different proposition.

So the idea of writing a book has always lived somewhere in the back of my mind.

Not because I thought I had nothing to say.

Almost the opposite.

I had too much to say, and the distance between the ideas in my head and a finished manuscript felt enormous.

Then something changed.

AI arrived.

Not AI as we had understood it for decades. Not machine learning quietly predicting which customer might leave or which transaction might be fraudulent. I mean the current generation of artificial intelligence: systems we can converse with, challenge,

question, collaborate with and increasingly ask to perform work alongside us.

I have spent much of my professional life working with this technology. And the more capable it became, the more I began to notice something profoundly personal.

AI was removing friction.

Things I could think but found difficult to express in writing became easier to communicate.

A chaotic collection of ideas could become structured.

A conversation could become an argument.

A brain dump could become prose.

The thoughts were still mine. The experiences were mine. The opinions were mine. The questions were mine.

But suddenly I had something capable of helping me bridge the gap between thinking and execution.

For perhaps the first time, I stopped seeing my dyslexia principally in terms of the limitations I had spent a career working around.

I started thinking differently about the *different thinking* that came with it.

AI could help with the gaps.

That meant I could concentrate on the part I was good at.

Thinking.

Questioning.

Connecting.

Being curious.

And that is how this book came into existence.

This book was not written in the normal way

I did not sit at a desk every evening and type chapter after chapter into a word processor.

In many cases, I didn't write the first draft at all.

I talked.

We began with a thesis.

As AI becomes increasingly capable of operating the tools we use to do our jobs, something fundamental happens to the value of knowing how to operate those tools ourselves.

For decades, we have built careers around proficiency.

We have hired people because they know Salesforce, SAP, Excel, Photoshop, Python, a particular cloud platform or whatever collection of tools happens to define their profession.

We have trained people in them.

Certified people in them.

Promoted people because of their mastery of them.

People have built identities around being exceptionally good at using them.

But what happens when AI can operate the tool?

If I can describe an outcome and an AI system can increasingly work out which buttons need pressing, which code needs writing, which application needs opening and which process needs following, where does human value move?

That question became an obsession.

And it led to another:

What if we have been measuring the wrong thing?

Perhaps knowing how to operate the tool was never the valuable part.

Perhaps it was simply the best proxy we had for something much harder to measure.

Judgement.

Experience.

Imagination.

Context.

The ability to recognise that something isn't quite right.

The ability to see connections that somebody else misses.

And, above all, the ability to ask a better question.

That became the idea behind what I call the **Curiosity Economy**.

An economy in which the scarcity increasingly isn't our ability to execute an idea.

It is our ability to have one.

That is an enormous change.

And once I became curious about it, I did what I have always done.

I started asking questions.

Curiosity created this book

The process we used to write the book you are holding is deliberately important.

I created the argument and structure I wanted to explore. Then, rather than asking me to write each chapter, AI interviewed me.

It asked questions.

I answered them.

Sometimes I answered coherently.

Sometimes I rambled.

Sometimes I started answering one question and ended up somewhere completely different because halfway through I realised there was another question that was more interesting.

I told stories from my career.

I described things I had seen inside organisations.

I challenged assumptions.

I contradicted myself.

I developed ideas as I spoke.

And when something didn't quite stand up, the AI could challenge me again.

Why do you believe that?

What's the evidence?

Doesn't that contradict what you said earlier?

What would somebody who disagrees with you say?

What happens if we take that argument to its logical conclusion?

Those questions mattered.

Because the AI's job wasn't to replace my curiosity.

Its job was to help me do something with it.

Then came the division of labour that would have been almost impossible for me before.

I supplied the curiosity, experiences, opinions, stories, arguments and questions.

AI helped turn them into the words on these pages.

That distinction is important.

I haven't outsourced my thinking.

I have outsourced some of the friction involved in expressing it.

And there is a world of difference between the two.

In fact, I think that difference is one of the most important things we need to understand about AI.

The most interesting future isn't one in which AI thinks so that humans don't have to.

It is one in which AI dramatically increases what humans can do with their thinking.

This book is evidence of that.

The organisation is no longer limited by technology

For most of modern economic history, ideas have been constrained by execution.

You might have a fantastic idea for a piece of software, but if you couldn't code and couldn't afford developers, it probably remained an idea.

You might have wanted to analyse a huge dataset, but if you didn't understand the analytical tools, you couldn't do very much with it.

You might have wanted to create a film, build a company, design a product, conduct research or write a book.

But between curiosity and outcome stood an enormous wall of specialist capability.

So we learned to value the people who could operate the machinery.

That made perfect sense.

But AI is beginning to lower that wall.

It doesn't eliminate expertise. It doesn't make experience irrelevant. And it certainly doesn't mean that everyone can suddenly do everything brilliantly.

But it changes the starting point.

Increasingly, the question is becoming less:

Do you know how to do this?

And more:

Do you know what is worth doing?

That shift has enormous implications for individuals.

It has even bigger implications for organisations.

I have come to believe that **the organisation is no longer limited by technology. It is limited by the questions it asks.**

Think about what that means.

For years, organisations have accumulated enormous backlogs of things they would like to do.

Ideas that weren't economical.

Questions nobody had time to investigate.

Processes that weren't quite broken enough to justify fixing.

Experiments that could never make it to the top of the priority list.

Small opportunities buried beneath bigger ones.

Technology created a scarcity of execution.

AI begins to attack that scarcity.

And if execution becomes dramatically cheaper and more abundant, another constraint has to take its place.

I believe that constraint is curiosity.

Who notices the opportunity?

Who asks why?

Who challenges the assumption everyone else has accepted?

Who connects something happening in one industry with a problem in another?

Who asks the question that wasn't in the project plan?

Who looks at something working perfectly well and wonders whether it could work completely differently?

Those behaviours have always been valuable.

We just haven't always valued them.

Now I think we are going to have to.

This isn't another book about becoming an AI expert

There are already plenty of books explaining artificial intelligence.

This isn't one of them.

You don't need to understand neural networks to understand the argument I'm going to make.

You don't need to become a prompt engineer.

You don't need another certificate.

And you certainly don't need to spend the next decade desperately trying to stay one step ahead of AI by learning every new tool that appears.

I think that may be precisely the wrong response.

Instead, I want to explore what happens when we stop competing with machines on their terms and start investing more heavily in the characteristics that make humans valuable in the first place.

Curiosity sits at the centre of that.

Throughout this book, I am going to challenge some deeply embedded assumptions about work.

Why do our CVs look the way they do?

Why do job adverts still ask for arbitrary numbers of years using particular technologies?

Why do organisations spend enormous amounts of money training people to operate platforms that AI will increasingly operate for them?

Why do we say we want innovative employees and then build cultures that punish people for questioning established ways of working?

Why do we measure productivity through output while often ignoring the person who asked the question that made half the output unnecessary?

And perhaps most importantly:

What should we do instead?

I don't want this to be a philosophical argument that leaves you nodding enthusiastically and then returning to work on Monday morning exactly as before.

The Curiosity Economy should change how we hire.

How we train.

How we educate.

How we manage.

How we promote.

How we think about experience.

How we use AI.

And how we think about our own value.

Because there is another side to this argument that we need to confront.

For somebody who has spent fifteen years becoming exceptionally good at a particular technology, being told that technology can increasingly perform the thing you have built your career around does not feel liberating.

It feels threatening.

I understand that.

This book will not dismiss the expertise people have spent their lives developing. Nor will it pretend that technological transitions are painless.

But I will make the case that something potentially extraordinary sits on the other side.

The skills being commoditised may not have been the most valuable human skills after all.

They may simply have been the skills the previous economy was best able to measure.

You are already holding the experiment

There is a temptation with any book about the future to ask whether its predictions will turn out to be correct.

This one is slightly different.

Part of the experiment has already happened.

You're holding it.

A dyslexic technologist who spent years thinking that writing a book would probably remain somewhere at the back of his mind became curious about an idea.

I explored it with AI.

The AI questioned me.

I questioned it.

Those questions created more questions.

Ideas became conversations.

Conversations became arguments.

Arguments became chapters.

And eventually those chapters became this book.

That doesn't make the AI the author of my experiences.

It didn't live my career.

It didn't sit in the meetings I sat in.

It didn't make my mistakes.

It didn't develop my beliefs for me.

But equally, it would be dishonest to pretend it was simply a more sophisticated spellchecker.

It has been a collaborator in turning thought into form.

That is why I wanted to be completely transparent about how this book was created.

I don't think using AI diminishes the work.

I think hiding its contribution would.

More importantly, I think the process demonstrates the argument I am about to make.

AI gave me access to capabilities that previously carried a much higher personal cost.

And rather than making my own abilities less relevant, something unexpected happened.

It made the characteristics that were most distinctly mine *more* useful.

My experiences mattered more.

My judgement mattered more.

My connections mattered more.

My questions mattered more.

My curiosity mattered more.

That is the possibility I want to explore with you.

Because if AI can increasingly compensate for the things we cannot do, perhaps the question we should be asking isn't:

What will be left for humans when AI can do everything we can do?

Perhaps there is a much more interesting question.

What becomes possible for humans when our ability to execute is no longer the thing holding our curiosity back?

This book is my attempt to find out.

And, appropriately enough, it starts with a question.

The CV That Time Forgot

Back in the day, I had friends for whom certifications were everything.

And I really do mean everything.

There was always another certification to get. Another exam to sit. Another little logo that could be added to LinkedIn or the bottom of a CV. You'd get certified in one version of a platform, then a new version would come along and there would be another certification for that. And the more of them you collected, the more employable you became.

It worked as well.

Some of these people built incredibly successful careers doing it. They earned good money, got promoted and became recognised as experts because they knew a particular piece of technology inside out.

And at the time, that made complete sense.

If a company had spent millions putting a particular platform at the centre of its business, it needed people who knew how to use it. If you could prove that you knew exactly where to click, how to configure it, how to run a particular process or how to

make the system produce something that everybody else needed, you were valuable.

There was knowledge trapped inside the interface.

You knew something other people didn't.

I remember looking at people who accumulated these certifications and thinking there was almost a science to it. You could plot a career through them. Get this certification, qualify for this role. Spend three years doing that, get the next certification, move up another level. Learn the next platform and suddenly another collection of jobs becomes available to you.

Entire careers were built around knowing how to make software do things.

There is absolutely nothing wrong with that. It was a rational response to the economy we had built.

But I've increasingly found myself wondering whether we ever stopped to ask what exactly we were valuing.

Because when you strip a lot of it back, what were we actually paying for?

Quite often, we were paying someone because they knew how to click around a screen better than somebody else.

That sounds deliberately reductive, and perhaps it is, but stay with me.

A significant proportion of what we call technical expertise has historically involved learning a sequence of interactions with a machine. You know which menu to open. You know which button to press. You know which field needs which information. You remember the keyboard shortcut. You know that if you want to produce this particular report, you go here, select this, change that, export it there and then do something else to it.

Eventually you become really, really good at it.

And because you're really good at it, you're faster than everybody else.

So we call you an expert.

Again, I'm not saying there isn't expertise surrounding that process. Of course there is. The best people understand why they're doing it, what the information means and what to do when something goes wrong.

But somewhere along the way we allowed ability to operate the tool to become a proxy for ability to do the job.

You can see it in almost every job description.

Think about the way they're written.

Right at the top are all the things you're expected to know how to do.

Five years' experience with this.

Three years' experience with that.

Must understand this platform.

Must have worked with this methodology.

Ideally certified in whatever technology happens to be important to the company writing the advert.

And then you get towards the bottom.

Must be a good communicator.

Must work well with others.

Should be creative.

Able to interpret information.

Comfortable dealing with ambiguity.

Curious.

They're almost written as the nice-to-haves.

And I think we're about to flip the whole bloody thing upside down.

Because AI is going to become better at operating those tools than we are.

Actually, that's probably understating it.

AI has the potential to become the greatest super-user we've ever created.

We have spent years teaching humans how to navigate graphical user interfaces. AI doesn't have to care about the graphical user interface. It can understand the interface, but increasingly it can also understand what's underneath it. It can work with the command line. It can call the API. It can understand the documentation. It can write the code. It can understand what individual functions actually do.

Most importantly, it can take what somebody is trying to achieve and work out how to translate that intention into actions across those systems.

That is the bit I think we are massively underestimating.

For most of my working life there has been this gap between wanting something to happen and knowing how to make the technology make it happen.

That gap created enormous amounts of economic value.

You wanted some information out of the business, so you found the person who knew how to run the report.

You wanted something changed in a system, so you found the person who understood the system.

You had an idea for an application, so you found somebody who could code.

You wanted to understand the data, so you found somebody who knew the analytical platform.

Often, success wasn't just about what you knew. It was about who you knew who knew how to do something.

I think most people who have worked in a large organisation will recognise this.

There is always that person.

"Speak to Sarah. She knows how to get that out of the system."

Or:

"Ask James. He built the original report."

Or:

"I'm sure Finance can do it, but I don't know who you need to speak to."

The organisation's capability is fragmented across people who know how to operate different bits of its machinery.

And suddenly we're building technology capable of operating the machinery for us.

That changes something much more fundamental than productivity.

It changes what the human is there to do.

Because if I don't need to know exactly how to operate the tool anymore, the important thing becomes what I can ask the tool to do.

And that's where this gets exciting.

I don't see this as some depressing race between humans and machines where we gradually discover all the things machines can do better than us until there's nothing left.

I see almost the opposite.

I think we've spent decades measuring human beings by some of the least interesting things about them.

How well can you remember this process?

How quickly can you navigate this interface?

Have you used version seven of this platform?

Can you reproduce the procedure somebody taught you three years ago?

And now we have a technology arriving that is phenomenally well suited to exactly those sorts of things.

Great.

Let it do them.

Because what becomes more interesting is what happens to all the things we stuck at the bottom of the job description.

Can you communicate an idea really well?

Can you take something that's half formed in your head and articulate it clearly enough for other people — and increasingly machines — to help you make it real?

Can you look at information and see something somebody else hasn't seen?

Can you connect an idea from one industry with a problem in another?

Can you work with other people?

Can you challenge an assumption?

Can you imagine something that doesn't exist yet?

Can you ask a really good question?

Those suddenly aren't the nice-to-haves.

They're the job.

And that, for me, is the really exciting bit.

Because it means that the next generation of work potentially values people for something much closer to what makes them uniquely interesting in the first place.

Two people might have access to exactly the same AI.

They might have access to exactly the same company data and exactly the same underlying tools.

One of them asks it to produce the same monthly report they've always produced.

The other asks, Why do we produce this report at all? What aren't we looking at? Is there another dataset we could combine with this? What would happen if we looked at the problem completely differently? What could we know that we've never previously been able to ask?

Same AI.

Same tools.

Same access.

Completely different value.

That's the shift.

When everybody has access to incredibly capable technology, knowing how to operate the technology stops being the thing that differentiates you.

What you can think to do with it starts to matter more.

And that opens the door to people who historically would have reached the edge of their technical ability and had to stop.

You might have had an amazing idea but couldn't code it.

You might have wanted to interrogate some data but didn't know SQL.

You might have had an idea for a business but couldn't build the website.

You might have wanted to research something but didn't know where to begin.

You might have been able to see exactly what was wrong with a process but had no idea how to change the underlying system.

So your curiosity had a ceiling.

Either you learnt the tool, paid somebody who had, or knew somebody who could help you.

AI starts removing that ceiling.

And when that happens, the question we ask about people changes.

For years we've effectively asked:

What do you know how to do?

I think we're moving towards something much more interesting:

What can you think of doing?

That's a fundamentally different employment market.

And it is why I think we're going to look back at some of today's job descriptions and find them quite strange.

We will wonder why we put so much emphasis on how many years somebody had spent operating a particular piece of

software, and so little emphasis on what they had actually thought about while they were operating it.

Because the human value was never really in clicking the button.

We just needed humans to click it.

Now that we increasingly don't, we get to ask a far more interesting question.

What happens when human curiosity is no longer limited by our ability to operate the tools?

That's the question this book is really about.

The Threshold Has Been Crossed

There was a time when I thought becoming better with technology meant learning more technology.

I have spent a large part of my career doing exactly that. New platforms arrived, new interfaces appeared, new programming languages became fashionable, and new certifications emerged to prove that you knew how to use all of them.

The logic was fairly straightforward. If I wanted to achieve something with a piece of software, I needed to understand how that software worked.

I needed to know where things were.

I needed to understand the menus. The terminology. The settings. The little quirks that were never quite captured in the documentation. I needed to know that if I put the wrong type of information into a particular field it would fail, or that some obscure option buried three screens deep would suddenly make the thing I was trying to achieve possible.

That knowledge was valuable.

In fact, as we explored in the previous chapter, for decades we built careers around it.

Then something changed.

And I don't think we have fully appreciated quite how fundamental that change is.

I experienced it personally through something surprisingly mundane: Canva.

I am not a designer. I have never pretended to be one. I can recognize things I like and things I don't like, and I can usually explain what I am trying to communicate, but put me in front of a blank design application and I am certainly not going to claim that I know exactly what to do.

I was trying to create some communications for Microsona, one of the products we had been developing at Imitation AI. I wanted different pieces of content. I wanted to experiment with messaging. I wanted imagery, layouts and punchy little statements that communicated what we were trying to say.

None of this was particularly extraordinary.

The interesting bit was how I did it.

Instead of opening Canva and deciding that I first needed to learn Canva, I started with the outcome.

I had already spent an enormous amount of time talking to AI about Microsona. It understood the product. It understood the language I tended to use. It understood the ideas I was trying to communicate and the style I was developing around them.

So I essentially said: you know what we have been talking about. You understand Microsona. You understand the sort of message I am trying to convey. I am not a designer. Help me turn this into something.

And it did.

It interacted with Canva and started bringing those ideas to life.

It selected imagery. It helped develop short pieces of copy. It worked within existing design approaches and templates. Within minutes, I had something tangible in front of me.

Was it the world's greatest piece of graphic design?

Probably not.

That isn't the point.

The point is that something had disappeared from between my idea and the outcome.

The tool.

Or, more accurately, my need to become proficient in the tool.

That was the moment that really started to unlock something for me.

For most of the history of computing, there has been an interface sitting between human intention and machine capability.

We want something to happen. The computer is capable of making it happen. But we need to learn the language required to tell the computer what to do.

Sometimes that language is literally code.

Sometimes it is a command line.

Most of the time, particularly over the last few decades, it has been a graphical user interface.

Buttons. Menus. Forms. Drop-down boxes. Tabs.

We have become so accustomed to this that we barely notice how strange the arrangement actually is.

Think about a large enterprise application.

A human expert might know ten screens extraordinarily well. Perhaps twenty. Someone who has spent years working with the

platform might know considerably more. They understand the sequence of actions required to complete a task. They know what tends to go wrong. They know which fields reject certain information. They know the workarounds.

But almost nobody knows one hundred percent of a sufficiently complicated application.

It is simply too big.

There are features they have never needed. Settings they have never touched. Functions hidden in parts of the platform belonging to somebody else's job.

And even the best human operator has another rather inconvenient limitation.

There is only one of them.

They can only do so much at once.

AI changes both of those things.

Give an AI system access to the documentation, tools and interfaces it needs and it can develop an extraordinarily broad understanding of how that software operates. Increasingly, it doesn't need to interact with the software in the same way we do at all. It can work beneath the interface, interacting directly with the capabilities the interface was designed to expose to us.

The UI was built for humans.

AI doesn't necessarily need it.

That distinction matters enormously.

Because once AI can operate the software, my relationship with that software changes.

I don't necessarily need to tell it:

Open this menu.

Select this option.

Copy that value.

Paste it here.

Run this report.

Export that file.

Open another application.

Import the file.

Change this setting.

Send the result.

Increasingly, I can describe what I actually want.

And that is a completely different relationship with technology.

After Canva, I started noticing the same pattern everywhere.

Email templates. Cloud infrastructure. Code. Communications. Development environments. Business processes.

Initially, AI might do one part for me.

Then two.

Then three.

Eventually, I would find myself looking at an entire workflow and realizing that much of the mechanical execution had disappeared from my workload.

The transition was incremental enough that it would have been easy to miss.

But once I saw it, I couldn't unsee it.

The important skill was moving.

It was no longer primarily about becoming better at operating the tool.

It was becoming better at explaining what I wanted from it.

The question was becoming the work.

That sentence is much more important than it might initially sound.

Because there is a temptation, particularly at the moment, to look at AI and ask:

What can I make this thing do?

I see this constantly.

People discover that an AI system can create a website, so they create a website.

They discover it can generate a video, so they generate a video.

They discover it can build an application, so they build an application.

They discover it can automate a process, so they automate a process.

And then we collectively complain about the enormous amount of AI-generated slop appearing around us.

But the problem isn't necessarily that AI produces slop.

Often the problem is that we haven't really decided what we want before asking it to produce something.

We are starting with capability rather than curiosity.

Look what this can do.

Rather than:

What am I actually trying to achieve?

That distinction is going to become increasingly important.

When execution was expensive, the friction of execution forced a certain amount of thinking.

If building a piece of software required a team of engineers, a budget, six months of development and organizational approval, somebody normally had to make a reasonably convincing case for why it should exist.

Not always, admittedly. Anyone who has worked in a large organization will know there are spectacular exceptions.

But friction acted as a filter.

AI is rapidly removing that friction.

And when the cost of making something collapses, the quality of deciding what should be made becomes far more important.

That means we have to move some of our discipline upstream.

Before building the solution, understand the problem.

Before asking AI to create the workflow, understand the outcome.

What does good look like?

Who is this actually for?

What happens when something goes wrong?

What constraints exist?

What assumptions are we making?

What haven't we thought about?

What would make this genuinely useful rather than merely impressive?

Ironically, AI itself can help us do this.

Rather than immediately asking it to build something, ask it to interrogate the idea.

Challenge me.

What have I missed?

What assumptions am I making?

What would stop this working?

What would a user expect that I haven't considered?

What questions should I be able to answer before we build anything?

That interaction is far more valuable than simply typing a vague request and accepting whatever comes back.

And this is where curiosity starts to become economically significant.

Because if execution becomes easier, the value moves toward the quality of the intention that precedes it.

I experienced this much more dramatically when I moved beyond using AI to help me write bits of code and started using it to help me deploy complete architectures.

I am not a full-stack engineer.

There are people who have spent their careers becoming exceptional full-stack engineers, and I would never claim that a few conversations with an AI magically transfer their experience into my head.

But something interesting happened.

I discovered I didn't necessarily need to become one.

I could give AI that role.

That is a very different claim.

I could explain what I was trying to build. I could question the architecture. I could understand the components. I could challenge decisions. I could test the outputs. I could ask why something had been implemented in a particular way.

But I didn't personally have to perform every technical action required to turn the idea into a working system.

And before somebody understandably says, "Yes, but AI-generated code can be insecure," I agree.

It can.

AI makes mistakes.

It misunderstands requirements. It makes assumptions. It can produce code that works but isn't necessarily the way an experienced engineer would have approached the problem.

So I tested it.

In one example, I was deploying into Google Cloud. I used Google's own security capabilities to assess the infrastructure. I ran security scans. I tested the authentication. The work subsequently went through ISO 27001-related scrutiny and penetration testing.

That distinction is critical to understanding what I am arguing.

I didn't say:

"AI built it, therefore it must be right."

I said:

"AI built it. Now how do we establish whether it is right?"

Those are profoundly different mindsets.

Working effectively with AI requires a relationship much closer to working with another person than operating a traditional piece of software.

If somebody new joins your team, you don't normally give them one instruction on Monday morning and expect flawless performance forever.

You explain.

They attempt something.

You review it.

Something isn't quite right.

You correct it.

They learn.

Next time it gets better.

The graduate working with you didn't know everything on their first day either.

Neither did you.

Yet there is a peculiar standard being applied to AI at the moment where a person asks it to do something once, spots an error and declares the entire technology useless.

I don't work with it like that.

When something is wrong, I correct it.

But importantly, I don't just correct the individual output. I feed the learning back into the way the system works.

This went wrong.

Here is why.

Next time, do this.

Then I watch.

Over time the mistakes reduce. The quality improves. More things come back right first time.

That feedback loop is where much of the productivity gain actually appears.

And eventually you notice something strange.

The AI isn't simply helping you complete your work.

It is changing what you consider your work to be.

I don't want to spend my day building every process inside my business.

I want to spend my time thinking about whether we are building the right business.

I want to think about customers.

Products.

Ideas.

Problems.

Opportunities.

I want to persuade people to use the thing we have created rather than spend all of my time operating the machinery required to create it.

AI has allowed me to move some of my attention from building the business toward driving the business. And this is much more valuable

Where Does the Human Go?

There is a version of the AI story that starts with a fairly grim question: if the machines can do more of the work, what happens to the people?

I understand why that question comes first. We have spent generations building organizations around the assumption that human execution is scarce. If you wanted something done, somebody had to do it. Somebody had to operate the system, prepare the report, move the information, check the numbers, update the record, create the proposal or coordinate the process. So when AI starts taking on more of that execution, the instinctive response is to look at the people doing it and ask whether we still need them.

I think that is the wrong place to start.

The more interesting question is what becomes possible when execution stops being the limiting factor. Because if AI can take on more of the operational weight that sits underneath an organization, the opportunity is not simply to make the same company cheaper. It is to change the shape of the company itself.

Picture a fairly ordinary business with ten salespeople. Those ten people are out in the market trying to win customers. They

are talking to people, understanding needs, shaping propositions and persuading someone that what the company offers is worth buying. Their work is very obviously connected to the direction and growth of the business. If they sell more, the company grows. If they understand customers better, the company learns. If they spot a new opportunity, the company has somewhere new to go.

But those ten people do not operate alone. Behind them is a much larger supporting cast. A customer conversation might lead to a proposal, which needs pricing, which needs margin checked, which needs a contract, which needs information entered into systems, which eventually needs delivery, invoicing and reporting. None of that work is pointless. In the organization we have today, it is essential. The sale only becomes real because a lot of other people help turn it into something the business can actually deliver.

That is the pattern I want to focus on. Most organizations have a relatively small number of people spending a large proportion of their time directly driving strategically important outcomes, and a much larger number of people whose work enables those outcomes to happen. The line is never perfectly clean, and I am not suggesting Finance, Operations or any other supporting function is somehow less important. I am talking about where human attention is consumed. A surprising amount of organizational capacity is spent operating the machinery around value creation rather than creating new value itself.

This is where the phrase “strategic operator” matters. I do not mean strategist. I am not imagining an organization in which everybody suddenly sits around theorizing about markets or producing more strategy decks. A strategic operator is someone doing something strategically important: selling, designing a

better product, improving a service, creating a new proposition, solving a customer problem or finding an opportunity the organization did not previously see. These are practical acts. They move the business.

Historically, we could only have so many people spending most of their time that way because strategic activity creates an execution tail behind it. The more the business tries to do, the more people it needs to support the doing. A new product creates processes. A new customer creates administration. A new market creates systems, controls and reporting. Complexity grows and, quite rationally, the organization grows around it.

That is why large organizations end up with so much internal machinery. We add systems to manage complexity, then people to operate the systems. We add processes to coordinate the work, then people to manage the processes. We add reporting to understand whether it is all working, then people to produce and interpret the reporting. None of this happened because leaders woke up wanting bureaucracy. It happened because execution was scarce and complexity had to be managed somehow.

For most of modern business history, every idea eventually collided with the same question: who is actually going to do this?

That question shapes more than workload. It shapes which ideas survive. A person can notice a brilliant opportunity, but if pursuing it requires data from one team, development effort from another, budget approval from a third and a six-month place in a transformation backlog, there is a good chance the idea will simply wait. Sometimes that is sensible prioritization. Sometimes it is just the cost of execution quietly killing curiosity before it has a chance to become useful.

AI changes the economics of that constraint.

We are already seeing the first, obvious version of this. A salesperson can use AI to prepare for a meeting, summarize the conversation and draft the follow-up. A finance team can automate parts of analysis and reporting. A product team can prototype much faster than it could a few years ago. These are genuine productivity gains, and organizations should take them.

But if that is where the thinking stops, we are still using AI to optimize the company we already have.

We ask, “How can I make my ten salespeople more productive?” Perhaps they become 30 percent more efficient. Perhaps 50 percent. That is useful. But once AI starts absorbing more of the execution that surrounds those ten people, I think a much more provocative question appears:

Why do we only have ten people selling?

I do not mean that everyone should literally have a sales target. I mean: why is strategically important work concentrated in such a small part of the organization?

Return to that sales team. Today, a good customer meeting creates a trail of operational activity behind it. In an AI-enabled organization, more of that trail can happen through the technology. The meeting can update the CRM. The proposal can be assembled from the conversation and the organization’s rules. Pricing and relevant information can be drawn from the right systems without the salesperson needing to know which five applications sit underneath the process. The human increasingly describes the outcome; AI works across the machinery required to make the outcome happen.

That is the threshold we explored in the previous chapter. AI begins to act as an intermediary between human intention and

the tools. Once that happens, the execution tail can become dramatically smaller.

Now the interesting part begins.

What happens to the people whose time was previously consumed by that execution?

The narrow answer is that the organization needs fewer of them. There will absolutely be companies that respond that way, and in some situations it will be rational. If a process genuinely needs less human effort, pretending otherwise does not help anyone.

But reduction is only one possible use of the capacity we have released. The much bigger opportunity is to move those people closer to value creation.

Imagine someone in Finance who has spent years looking at the economics of the customer base. Today, much of their week might be consumed producing and reconciling information. If AI takes away a meaningful part of that work, their experience does not disappear with the spreadsheet. They still understand which customers are profitable, where margins move and which patterns never quite made sense. What if more of their time can now be spent pursuing those patterns rather than merely reporting them?

Or take someone in Operations who has spent years seeing where customers repeatedly struggle. They have probably developed instincts about what is broken long before it appears in a strategic plan. In the old organization, turning that instinct into action might have required a project, a team and months of coordination. In the new one, AI can help them investigate the

pattern, model an alternative and take the idea far enough that it becomes something the organization can actually evaluate.

That is the opportunity I mean when I talk about widening strategic capacity. We are not making a small number of strategic operators slightly better. We are enabling far more people to do strategically important things.

And that is a radically different ambition.

Most organizations today have narrow strategic capacity. A relatively small proportion of the workforce has enough time, authority and access to pursue new customers, test new ideas or redesign how something works. The rest of the organization is not necessarily short of intelligence, experience or imagination. Often it is simply busy. People have a job to do, systems to update and processes to complete. Their curiosity has to fit around the work.

AI creates the possibility of reversing that relationship. Instead of curiosity fitting around execution, execution starts to fit around curiosity.

What would a company look like if 90 percent of its people could spend most of their time as strategic operators?

Again, not strategists. Operators.

They would still have different disciplines. The finance expert would still understand finance better than the salesperson. The person in Operations would still know the process in a way somebody in Product does not. That difference is valuable. The point is not to flatten expertise; it is to free expertise from being consumed by the mechanics of its own tools.

That distinction matters because the experience inside an organization is much richer than the roles we normally allow

people to play. We hire whole human beings and then often use very narrow slices of them. Someone might have fifteen years of customer experience, deep knowledge of a sector and a strong instinct for what is changing, yet spend a large part of their day administering a platform because that is what the role requires. When AI can operate more of the platform, the question becomes whether we can finally use more of the person.

That is why I find the productivity narrative around AI too small. Productivity asks how much faster the existing job can be done. I am interested in what the person can do once the existing job no longer consumes them in the same way.

If we get that right, the organization does not simply become more efficient. It becomes strategically wider. It has more people capable of finding customers, improving products and discovering opportunities. It can pursue more ideas because the cost of moving from idea to action has fallen. It can respond to more signals because the people closest to those signals can do something with them.

That last point is particularly important.

In a traditional organization, curiosity is often separated from execution by distance. Somebody at the frontline notices something. They tell a manager. The manager takes it to a meeting. It might enter a backlog, then compete for attention with everything else. By the time the organization is ready to act, the original observation may have lost its urgency or the person who noticed it may no longer be involved.

The curiosity existed. What it lacked was execution.

Now imagine the same person with AI available as a practical intermediary. They notice a pattern in customer behavior and can

immediately start investigating it. They can explore the available data, ask whether the pattern is real, test a few explanations and create a rough version of what a different approach might look like. They still need governance, judgement and the right permissions. AI does not remove those responsibilities. What it removes is some of the distance between noticing something and being able to do something useful about it.

That distance is where a huge amount of organizational curiosity currently dies.

When it collapses for one employee, that is useful. When it collapses for thousands, it changes the operating model of the company.

This is where curiosity stops being a soft skill and becomes an economic one.

If execution becomes abundant, the constraint moves upstream. The organization can create more, analyze more and automate more than it could before. But abundant capability does not guarantee valuable outcomes. Somebody still has to decide what is worth doing. Somebody has to notice the customer behavior worth investigating or the assumption worth challenging. Somebody has to ask the question that sends all that new execution power in a useful direction.

That is why I keep coming back to the same line:

The organization is no longer limited by technology. It is limited by the questions it asks.

In the old world, a great question could still be economically useless if acting on it was too expensive. In the new world, the cost of acting on a question is falling quickly. That makes the quality of the question itself more valuable.

This is also why I do not believe the future organization is simply the current organization with fewer people. That may be the first-order response because it is easy to measure. If AI saves ten thousand hours, a spreadsheet can convert those hours into headcount. The alternative is harder because it asks leaders to imagine what the organization could do with ten thousand hours of newly available human capacity.

Could those hours create new revenue? Could they improve retention? Could they uncover a product the company has not thought of? Could they allow people who understand the customer in completely different ways to act on what they know?

Those questions are harder to put into a cost-reduction business case, but strategically they are far more interesting.

There is a competitive dimension here as well. Imagine two companies adopt broadly similar AI capability. The first uses it primarily to automate existing work and reduce cost. The second takes the same efficiency, but deliberately reinvests a large part of it into widening strategic capacity. It gives more people permission and tools to explore customer needs, improve products and pursue ideas that previously would have died in the queue.

At first, the first company may look more efficient. But over time I suspect the second becomes much harder to compete with, because it has multiplied the number of people who can drive the business.

This shift also changes the role of leadership. If far more people can act, the answer cannot simply be “everyone go and be curious.” Ten thousand disconnected experiments are not a strategy. Curiosity without direction becomes noise.

The job of leadership therefore becomes creating a strong enough sense of direction that curiosity can cascade through the organization. Leaders still decide what the organization is trying to achieve, which outcomes matter and where its boundaries are. But instead of translating that direction into ever more detailed instructions, they create questions that people throughout the company can pursue from their own context.

Suppose the organization wants to grow in a particular market. The salesperson closest to that market may ask which customers are being ignored. Someone in Operations may see a service constraint that would prevent growth. A product expert may understand which part of the proposition needs to change. They are not all trying to do each other's jobs. They are contributing different pieces of curiosity to the same strategic outcome.

AI allows each of them to take their piece further.

That is how curiosity can cascade without becoming an overhead. Direction comes from the organization; questions multiply through the people; AI helps turn those questions into investigation and execution; the learning flows back into the organization. The cycle then starts again with better information.

This is a very different rhythm from the organization in which thought happens at the top and execution happens underneath it.

It is also why the phrase “strategic operator” matters so much to me. I do not want to replace a hierarchy of executors with a hierarchy of amateur strategists. I want to expand the number of people who can do work that changes the outcome of the business.

That might mean selling. It might mean designing a better product. It might mean solving a customer problem that has

quietly existed for years. The exact act will differ from person to person, but the common feature is that the person is moving the organization rather than simply moving work through the organization.

Once you see that distinction, the question “Where does the human go?” starts to look different.

The human moves closer to the outcome.

AI increasingly handles the distance between intention and execution. The human spends more time deciding which intentions are worth pursuing, bringing judgement, context, experience and curiosity to the moment before the execution begins.

That is not the disappearance of human value. It is a chance to apply human value much more widely.

Today, most organizations have a relatively small group of strategic operators supported by a much larger execution machine. Tomorrow, AI gives us the possibility of compressing that execution layer and redirecting a much greater proportion of human capacity toward the work that actually moves the business.

Not ten strategic operators supported by twenty executors, but thirty people with the ability to operate strategically. Then hundreds. Then thousands.

Perhaps, eventually, almost everybody.

That is the enterprise arc I believe AI makes possible. And curiosity is what unlocks it, because when execution becomes abundant, the scarce capability is no longer knowing how to make the machinery work. It is knowing what is worth asking the machinery to do.